

HIGH-PERFORMANCE COMPUTING

Stepan Gordeev

September 17, 2020

OUTLINE

1. Parallelization
2. BlueHive
3. General Coding
4. Julia Recommendations
5. Exposition of Advanced Techniques

PARALLELIZATION

OVERVIEW

- All consumer CPUs have multiple cores, but by default your code will run sequentially on one core!
- Can write code that will explicitly split tasks across cores.
- Work is split between cores, they all work simultaneously, then they pool results together.
- Matlab and Julia will run certain things in parallel on their own, e.g. linear algebra.
- Split sequential code across two processors \rightarrow 2x speedup
 - Parallelization adds overhead: CPU does extra work to split the task across cores and then to collect the results.

OVERVIEW

- All consumer CPUs have multiple cores, but by default your code will run sequentially on one core!
- Can write code that will explicitly split tasks across cores.
- Work is split between cores, they all work simultaneously, then they pool results together.
- Matlab and Julia will run certain things in parallel on their own, e.g. linear algebra.
- Split sequential code across two processors \rightarrow 2x speedup
 - Parallelization adds overhead: CPU does extra work to split the task across cores and then to collect the results.

OVERVIEW

- All consumer CPUs have multiple cores, but by default your code will run sequentially on one core!
- Can write code that will explicitly split tasks across cores.
- Work is split between cores, they all work simultaneously, then they pool results together.
- Matlab and Julia will run certain things in parallel on their own, e.g. linear algebra.
- Split sequential code across two processors \rightarrow 2x speedup
 - Parallelization adds overhead: CPU does extra work to split the task across cores and then to collect the results.

OVERVIEW

- All consumer CPUs have multiple cores, but by default your code will run sequentially on one core!
- Can write code that will explicitly split tasks across cores.
- Work is split between cores, they all work simultaneously, then they pool results together.
- Matlab and Julia will run certain things in parallel on their own, e.g. linear algebra.
- Split sequential code across two processors \rightarrow 2x speedup
 - Parallelization adds overhead: CPU does extra work to split the task across cores and then to collect the results.

OVERVIEW

- All consumer CPUs have multiple cores, but by default your code will run sequentially on one core!
- Can write code that will explicitly split tasks across cores.
- Work is split between cores, they all work simultaneously, then they pool results together.
- Matlab and Julia will run certain things in parallel on their own, e.g. linear algebra.
- Split sequential code across two processors \rightarrow 2x speedup
 - Parallelization adds overhead: CPU does extra work to split the task across cores and then to collect the results.

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHEN TO PARALLELIZE

- Cannot parallelize tasks that have to be done sequentially
 - Outer loop of VFI (go through V_l): No (V_{l+1} depends on V_l).
 - Inner loop of VFI (go through k_i): Yes (k_{i+1} does not depend on k_i). Send the $V_{l+1}(k_i) = \max \dots$ problem to each core, then collect results before moving on to V_{l+2} .
- Run $V_{l+1}(k_i) = \max \dots$ tasks in parallel:
 - e.g. CPU has 4 cores, grid is $[k_1, \dots, k_{100}]$
 - core 1 solves $V_{l+1}(k_i) = \max \dots$ for $i = 1, \dots, 25$
 - core 2 solves $V_{l+1}(k_i) = \max \dots$ for $i = 26, \dots, 50$
 - ...
 - CPU collects four pairs of $V'(k)$, $g(k)$, merges into one pair
 - Move on to V_{l+2}

WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
 for j=1:100
 ...
 end
end
```
- if parallelize outer *i* loop: 100 tasks, each has 1/100th work
  - CPU sends out work to cores and collects results just once
- if parallelize inner *j* loop: 10,000 tasks, each has 1/10,000th work
  - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

## WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
    for j=1:100
        ...
    end
end
```
- if parallelize outer *i* loop: 100 tasks, each has 1/100th work
 - CPU sends out work to cores and collects results just once
- if parallelize inner *j* loop: 10,000 tasks, each has 1/10,000th work
 - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
 for j=1:100
 ...
 end
end
```
- if parallelize outer **i** loop: 100 tasks, each has 1/100th work
  - CPU sends out work to cores and collects results just once
- if parallelize inner **j** loop: 10,000 tasks, each has 1/10,000th work
  - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

## WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
    for j=1:100
        ...
    end
end
```
- if parallelize outer **i** loop: 100 tasks, each has 1/100th work
 - CPU sends out work to cores and collects results just once
- if parallelize inner **j** loop: 10,000 tasks, each has 1/10,000th work
 - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
 for j=1:100
 ...
 end
end
```
- if parallelize outer **i** loop: 100 tasks, each has 1/100th work
  - CPU sends out work to cores and collects results just once
- if parallelize inner **j** loop: 10,000 tasks, each has 1/10,000th work
  - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

## WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
    for j=1:100
        ...
    end
end
```
- if parallelize outer **i** loop: 100 tasks, each has 1/100th work
 - CPU sends out work to cores and collects results just once
- if parallelize inner **j** loop: 10,000 tasks, each has 1/10,000th work
 - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

WHAT TO PARALLELIZE

- Split tasks to maximize computation by cores and minimize communication between cores
- ```
for i=1:100
 for j=1:100
 ...
 end
end
```
- if parallelize outer **i** loop: 100 tasks, each has 1/100th work
  - CPU sends out work to cores and collects results just once
- if parallelize inner **j** loop: 10,000 tasks, each has 1/10,000th work
  - CPU sends out work to cores and collects results 100 times
- outer is faster: less overhead

## HOW TO PARALLELIZE

- Julia: several ways, threads are simplest.
- Julia needs to be started with a pre-set number of threads.
  - Juno will pass this setting automatically.
  - Explicitly: Preferences → Packages → Julia Client → Number of Threads
- Verify threads setting: `Threads.nthreads()`
- Put `Threads.@threads` before a loop you want to make parallel:  

```
Threads.@threads for i=1:n
 ...
end
```
- In Matlab, use `parpool` to control the number of worker processes and replace `for` with `parfor`.

## HOW TO PARALLELIZE

- Julia: several ways, threads are simplest.
- Julia needs to be started with a pre-set number of threads.
  - Juno will pass this setting automatically.
  - Explicitly: Preferences → Packages → Julia Client → Number of Threads
- Verify threads setting: `Threads.nthreads()`
- Put `Threads.@threads` before a loop you want to make parallel:

```
Threads.@threads for i=1:n
 ...
end
```
- In Matlab, use `parpool` to control the number of worker processes and replace `for` with `parfor`.

## HOW TO PARALLELIZE

- Julia: several ways, threads are simplest.
- Julia needs to be started with a pre-set number of threads.
  - Juno will pass this setting automatically.
  - Explicitly: Preferences → Packages → Julia Client → Number of Threads
- Verify threads setting: **Threads.nthreads()**
- Put `Threads.@threads` before a loop you want to make parallel:

```
Threads.@threads for i=1:n
 ...
end
```
- In Matlab, use `parpool` to control the number of worker processes and replace `for` with `parfor`.

## HOW TO PARALLELIZE

- Julia: several ways, threads are simplest.
- Julia needs to be started with a pre-set number of threads.
  - Juno will pass this setting automatically.
  - Explicitly: Preferences → Packages → Julia Client → Number of Threads
- Verify threads setting: `Threads.nthreads()`
- Put `Threads.@threads` before a loop you want to make parallel:  

```
Threads.@threads for i=1:n
 ...
end
```
- In Matlab, use `parpool` to control the number of worker processes and replace `for` with `parfor`.

## HOW TO PARALLELIZE

- Julia: several ways, threads are simplest.
- Julia needs to be started with a pre-set number of threads.
  - Juno will pass this setting automatically.
  - Explicitly: Preferences → Packages → Julia Client → Number of Threads
- Verify threads setting: `Threads.nthreads()`
- Put `Threads.@threads` before a loop you want to make parallel:  

```
Threads.@threads for i=1:n
 ...
end
```
- In Matlab, use `parpool` to control the number of worker processes and replace `for` with `parfor`.

BLUEHIVE

---

## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers



## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers

## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers

## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers

## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers

## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers

## OVERVIEW

- BlueHive is a university cluster running RHEL (a Linux distribution)
- Hundreds of compute nodes, each with 12-64 processors.
- Info: `info.circ.rochester.edu`
- Uses of BlueHive
  - Large parallel problems: exploit huge number of cores
  - Work with large datasets: exploit huge RAM
  - Expensive specialized software unavailable on lab computers

## ACCESSING BLUEHIVE

- Access through **bluehive.circ.rochester.edu**
  - if off campus, have to connect through the university VPN
- Launch session with desired parameters:  
`mate-session --time=8:00:00 --cpus-per-task=16 --mem=32g`
- Wait for the session to appear. Wait time depends on queue, faster if ask for fewer resources.
- Session opens in a virtual desktop. The desktop environment is MATE.
- Pop-up bar in the top: toggle fullscreen, exchange text between your computer's clipboard and the virtual desktop's clipboard.

## ACCESSING BLUEHIVE

- Access through `bluehive.circ.rochester.edu`
  - if off campus, have to connect through the university VPN
- Launch session with desired parameters:  
`mate-session --time=8:00:00 --cpus-per-task=16 --mem=32g`
- Wait for the session to appear. Wait time depends on queue, faster if ask for fewer resources.
- Session opens in a virtual desktop. The desktop environment is MATE.
- Pop-up bar in the top: toggle fullscreen, exchange text between your computer's clipboard and the virtual desktop's clipboard.



## ACCESSING BLUEHIVE

- Access through `bluehive.circ.rochester.edu`
  - if off campus, have to connect through the university VPN
- Launch session with desired parameters:  
`mate-session --time=8:00:00 --cpus-per-task=16 --mem=32g`
- Wait for the session to appear. Wait time depends on queue, faster if ask for fewer resources.
- Session opens in a virtual desktop. The desktop environment is MATE.
- Pop-up bar in the top: toggle fullscreen, exchange text between your computer's clipboard and the virtual desktop's clipboard.

## ACCESSING BLUEHIVE

- Access through `bluehive.circ.rochester.edu`
  - if off campus, have to connect through the university VPN
- Launch session with desired parameters:  
`mate-session --time=8:00:00 --cpus-per-task=16 --mem=32g`
- Wait for the session to appear. Wait time depends on queue, faster if ask for fewer resources.
- Session opens in a virtual desktop. The desktop environment is MATE.
- Pop-up bar in the top: toggle fullscreen, exchange text between your computer's clipboard and the virtual desktop's clipboard.

## ACCESSING BLUEHIVE

- Access through `bluehive.circ.rochester.edu`
  - if off campus, have to connect through the university VPN
- Launch session with desired parameters:  
`mate-session --time=8:00:00 --cpus-per-task=16 --mem=32g`
- Wait for the session to appear. Wait time depends on queue, faster if ask for fewer resources.
- Session opens in a virtual desktop. The desktop environment is MATE.
- Pop-up bar in the top: toggle fullscreen, exchange text between your computer's clipboard and the virtual desktop's clipboard.

## SOFTWARE

- BH has an extensive list of free and paid software pre-installed (everything that lab computers have, and more): see **`info.circ.rochester.edu/#BlueHive/Software/`** for the list.
- Some programs available in Applications menu (e.g. MATLAB, Stata)
- If a program isn't there, need to activate it through the terminal
  - see the software list to find the appropriate package name for the program you want
  - open terminal through the applications menu or a top panel shortcut
  - type `module load package_name` and press Enter

## SOFTWARE

- BH has an extensive list of free and paid software pre-installed (everything that lab computers have, and more): see `info.circ.rochester.edu/#BlueHive/Software/` for the list.
- Some programs available in Applications menu (e.g. MATLAB, Stata)
- If a program isn't there, need to activate it through the terminal
  - see the software list to find the appropriate package name for the program you want
  - open terminal through the applications menu or a top panel shortcut
  - type `module load package_name` and press Enter

## SOFTWARE

- BH has an extensive list of free and paid software pre-installed (everything that lab computers have, and more): see `info.circ.rochester.edu/#BlueHive/Software/` for the list.
- Some programs available in Applications menu (e.g. MATLAB, Stata)
- If a program isn't there, need to activate it through the terminal
  - see the software list to find the appropriate package name for the program you want
  - open terminal through the applications menu or a top panel shortcut
  - type `module load package_name` and press Enter

## SOFTWARE

- BH has an extensive list of free and paid software pre-installed (everything that lab computers have, and more): see `info.circ.rochester.edu/#BlueHive/Software/` for the list.
- Some programs available in Applications menu (e.g. MATLAB, Stata)
- If a program isn't there, need to activate it through the terminal
  - see the software list to find the appropriate package name for the program you want
  - open terminal through the applications menu or a top panel shortcut
  - type `module load package_name` and press Enter

## SOFTWARE

- BH has an extensive list of free and paid software pre-installed (everything that lab computers have, and more): see `info.circ.rochester.edu/#BlueHive/Software/` for the list.
- Some programs available in Applications menu (e.g. MATLAB, Stata)
- If a program isn't there, need to activate it through the terminal
  - see the software list to find the appropriate package name for the program you want
  - open terminal through the applications menu or a top panel shortcut
  - type `module load package_name` and press Enter



## SOFTWARE

- BH has an extensive list of free and paid software pre-installed (everything that lab computers have, and more): see `info.circ.rochester.edu/#BlueHive/Software/` for the list.
- Some programs available in Applications menu (e.g. MATLAB, Stata)
- If a program isn't there, need to activate it through the terminal
  - see the software list to find the appropriate package name for the program you want
  - open terminal through the applications menu or a top panel shortcut
  - type `module load package_name` and press Enter

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
    - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - `load module, set threads`
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - `load module, set threads`
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.



## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the `uber-juno` package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the **uber-juno** package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the `julia-client` Atom package settings. Set the number of threads manually there as well.

## SOFTWARE: JULIA

- Running Julia shell
  - `module load julia`
  - `export JULIA_NUM_THREADS=16` (or however many CPUs you requested)
  - `julia`
  - you are now in the Julia REPL
- Running a .jl file
  - load module, set threads
  - `julia 'path/to/your/file.jl'`
- Using Atom as an IDE for Julia
  - `module load julia`
  - `module load atom`
  - `atom`
  - once Atom has launched, install the **uber-juno** package in it
  - find the path to the Julia executable by running `which julia`. Paste into the “Julia Path” setting of the **julia-client** Atom package settings. Set the number of threads manually there as well.

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page



## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page

## FILES AND JOBS

- Quick and dirty way to transfer files: upload to some cloud service, open its web app in Firefox on the virtual desktop, download.
- For more efficiency and speed, use FTP. See the info page.
- Jobs: don't need to run code through GUI
  - can set up a script asking for certain resources and executing a certain command (e.g. `julia file.jl`)
  - most software packages we use (MATLAB, Julia, Stata, R) can run scripts (.m, .jl, .do, .R) from the command line
  - the job enters a queue and is executed once resources are available
  - preferable for more demanding programs (don't need to wait around for session to open)
  - more efficient when want to run many scripts simultaneously
  - see the info page



## GENERAL CODING

---

## PRECISION

- Computers store approximations of real numbers in floating point type.
- Default: double precision (64 bits).
- In Julia: can go down to 16/32, can go up to 128 with packages, can do arbitrary precision.
- Can declare all variables as 32-bit types instead of 64-bit
  - RAM consumption ↓ almost 50%, speed ↑
  - (usually) negligible precision loss
- Machine epsilon for double:  $2^{-52} = 2.2 \times 10^{-16}$ .
- Results of intermediate computations are approximated:

```
julia> sqrt(2)^2 == 2
false
```

```
julia> isapprox(sqrt(2)^2, 2)
true
```

## PRECISION

- Computers store approximations of real numbers in floating point type.
- Default: double precision (64 bits).
- In Julia: can go down to 16/32, can go up to 128 with packages, can do arbitrary precision.
- Can declare all variables as 32-bit types instead of 64-bit
  - RAM consumption ↓ almost 50%, speed ↑
  - (usually) negligible precision loss
- Machine epsilon for double:  $2^{-52} = 2.2 \times 10^{-16}$ .
- Results of intermediate computations are approximated:

```
julia> sqrt(2)^2 == 2
false
```

```
julia> isapprox(sqrt(2)^2, 2)
true
```

## PRECISION

- Computers store approximations of real numbers in floating point type.
- Default: double precision (64 bits).
- In Julia: can go down to 16/32, can go up to 128 with packages, can do arbitrary precision.
- Can declare all variables as 32-bit types instead of 64-bit
  - RAM consumption ↓ almost 50%, speed ↑
  - (usually) negligible precision loss
- Machine epsilon for double:  $2^{-52} = 2.2 \times 10^{-16}$ .
- Results of intermediate computations are approximated:

```
julia> sqrt(2)^2 == 2
false
```

```
julia> isapprox(sqrt(2)^2, 2)
true
```

## PRECISION

- Computers store approximations of real numbers in floating point type.
- Default: double precision (64 bits).
- In Julia: can go down to 16/32, can go up to 128 with packages, can do arbitrary precision.
- Can declare all variables as 32-bit types instead of 64-bit
  - RAM consumption ↓ almost 50%, speed ↑
  - (usually) negligible precision loss
- Machine epsilon for double:  $2^{-52} = 2.2 \times 10^{-16}$ .
- Results of intermediate computations are approximated:

```
julia> sqrt(2)^2 == 2
false
```

```
julia> isapprox(sqrt(2)^2, 2)
true
```

## PRECISION

- Computers store approximations of real numbers in floating point type.
- Default: double precision (64 bits).
- In Julia: can go down to 16/32, can go up to 128 with packages, can do arbitrary precision.
- Can declare all variables as 32-bit types instead of 64-bit
  - RAM consumption ↓ almost 50%, speed ↑
  - (usually) negligible precision loss
- Machine epsilon for double:  $2^{-52} = 2.2 \times 10^{-16}$ .
- Results of intermediate computations are approximated:

```
julia> sqrt(2)^2 == 2
false
```

```
julia> isapprox(sqrt(2)^2, 2)
true
```

## PRECISION

- Computers store approximations of real numbers in floating point type.
- Default: double precision (64 bits).
- In Julia: can go down to 16/32, can go up to 128 with packages, can do arbitrary precision.
- Can declare all variables as 32-bit types instead of 64-bit
  - RAM consumption ↓ almost 50%, speed ↑
  - (usually) negligible precision loss
- Machine epsilon for double:  $2^{-52} = 2.2 \times 10^{-16}$ .
- Results of intermediate computations are approximated:

```
julia> sqrt(2)^2 == 2
false
```

```
julia> isapprox(sqrt(2)^2, 2)
true
```

## COMPARISON OF LANGUAGES

- Aruoba and Fernández-Villaverde (2014), “A Comparison of Programming Languages in Economics”:

| Language | Version/Compiler     | Time | Rel. Time |
|----------|----------------------|------|-----------|
| C++      | GCC-4.9.0            | 0.73 | 1.00      |
|          | Intel C++ 14.0.3     | 1.00 | 1.38      |
|          | Clang 5.1            | 1.00 | 1.38      |
| Fortran  | GCC-4.9.0            | 0.76 | 1.05      |
|          | Intel Fortran 14.0.3 | 0.95 | 1.30      |
| Java     | JDK8u5               | 1.95 | 2.69      |
| Julia    | 0.2.1                | 1.92 | 2.64      |
| Matlab   | 2014a                | 7.91 | 10.88     |

- Running time usually less important than development time.
- MATLAB and Julia run slower, but much faster to code and debug.
- Julia considerably faster than MATLAB



## COMPARISON OF LANGUAGES

- Aruoba and Fernández-Villaverde (2014), “A Comparison of Programming Languages in Economics”:

| Language | Version/Compiler     | Time | Rel. Time |
|----------|----------------------|------|-----------|
| C++      | GCC-4.9.0            | 0.73 | 1.00      |
|          | Intel C++ 14.0.3     | 1.00 | 1.38      |
|          | Clang 5.1            | 1.00 | 1.38      |
| Fortran  | GCC-4.9.0            | 0.76 | 1.05      |
|          | Intel Fortran 14.0.3 | 0.95 | 1.30      |
| Java     | JDK8u5               | 1.95 | 2.69      |
| Julia    | 0.2.1                | 1.92 | 2.64      |
| Matlab   | 2014a                | 7.91 | 10.88     |

- Running time usually less important than development time.
- MATLAB and Julia run slower, but much faster to code and debug.
- Julia considerably faster than MATLAB

## COMPARISON OF LANGUAGES

- Aruoba and Fernández-Villaverde (2014), “A Comparison of Programming Languages in Economics”:

| Language | Version/Compiler     | Time | Rel. Time |
|----------|----------------------|------|-----------|
| C++      | GCC-4.9.0            | 0.73 | 1.00      |
|          | Intel C++ 14.0.3     | 1.00 | 1.38      |
|          | Clang 5.1            | 1.00 | 1.38      |
| Fortran  | GCC-4.9.0            | 0.76 | 1.05      |
|          | Intel Fortran 14.0.3 | 0.95 | 1.30      |
| Java     | JDK8u5               | 1.95 | 2.69      |
| Julia    | 0.2.1                | 1.92 | 2.64      |
| Matlab   | 2014a                | 7.91 | 10.88     |

- Running time usually less important than development time.
- MATLAB and Julia run slower, but much faster to code and debug.
- Julia considerably faster than MATLAB

## COMPARISON OF LANGUAGES

- Aruoba and Fernández-Villaverde (2014), “A Comparison of Programming Languages in Economics”:

| Language | Version/Compiler     | Time | Rel. Time |
|----------|----------------------|------|-----------|
| C++      | GCC-4.9.0            | 0.73 | 1.00      |
|          | Intel C++ 14.0.3     | 1.00 | 1.38      |
|          | Clang 5.1            | 1.00 | 1.38      |
| Fortran  | GCC-4.9.0            | 0.76 | 1.05      |
|          | Intel Fortran 14.0.3 | 0.95 | 1.30      |
| Java     | JDK8u5               | 1.95 | 2.69      |
| Julia    | 0.2.1                | 1.92 | 2.64      |
| Matlab   | 2014a                | 7.91 | 10.88     |

- Running time usually less important than development time.
- MATLAB and Julia run slower, but much faster to code and debug.
- Julia considerably faster than MATLAB

# PROFILING

- Time and profile your code
  - Julia: `@time some_fn()` to measure time and memory allocations
  - Julia: Profile module to show time spent on each line (`@profile` and `@profiler`)
- Find bottlenecks, figure out how to optimize them

# PROFILING

- Time and profile your code
  - Julia: `@time some_fn()` to measure time and memory allocations
  - Julia: Profile module to show time spent on each line (`@profile` and `@profiler`)
- Find bottlenecks, figure out how to optimize them

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab



# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab






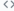

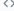

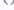

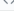

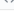

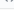

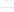
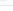
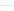
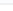
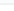
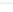
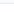
- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20ptimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20optimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab



- Version control: revert to *any* version of the code
  - backups and file copies are unreliable, inconvenient, and inflexible
  - a file name from my 2nd year project:  
`model2_4parallelVer20optimizedMore.jl`
  - instead, use a Version Control System
- Git: most popular VCS
  - change a file  $\implies$  *commit* the change with a descriptive name
  - Git will store this commit forever
  - can easily see differences between commits, revert to any commit
  - each commit takes up little space: only the difference from previous commit needs to be stored
- Collaboration
  - can work in parallel (*branches*), easily merge commits
  - host code on GitHub or GitLab

# GIT EXAMPLE: THIS COURSE

|                                                                                                                   |                                                                                                             |                                                                                     |
|-------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| Commits on Aug 30, 2020                                                                                           |                                                                                                             |                                                                                     |
| <b>Copied BH signup instructions to the notes</b><br>stepangordeev committed 3 days ago                           |  <a href="#">afd03df</a> |  |
| <b>Expanded BH instructions</b><br>stepangordeev committed 3 days ago                                             |  <a href="#">7a1d193</a> |  |
| <b>Second lecture VFI note made more explicit</b><br>stepangordeev committed 3 days ago                           |  <a href="#">b2006ad</a> |  |
| <b>Add uncovers to HPC lecture, expand on 32 bits</b><br>stepangordeev committed 3 days ago                       |  <a href="#">b086443</a> |  |
| <b>Dynare output slides</b><br>stepangordeev committed 3 days ago                                                 |  <a href="#">389d986</a> |  |
| <b>Merge branch 'hw3_dynare_calib'</b><br>stepangordeev committed 3 days ago                                      |  <a href="#">5c4761e</a> |  |
| <b>More detailed and explicit Dynare calibration instructions in hw3</b><br>stepangordeev committed 3 days ago    |  <a href="#">58ec23e</a> |  |
| <b>Merge branch 'hw3_dynare_calib'</b><br>stepangordeev committed 3 days ago                                      |  <a href="#">dad4f73</a> |  |
| <b>Adjusted hw3 dynare problem solution to correct parameter definition</b><br>stepangordeev committed 3 days ago |  <a href="#">c906f03</a> |  |
| <b>Greatly expanded Dynare estimation slides</b><br>stepangordeev committed 3 days ago                            |  <a href="#">7b81845</a> |  |
| <b>Clarified and expanded Dynare estimation slides</b><br>stepangordeev committed 3 days ago                      |  <a href="#">0951d71</a> |  |
| <b>Small Dynare lecture fixes</b><br>stepangordeev committed 3 days ago                                           |  <a href="#">0a2baf1</a> |  |

## GIT EXAMPLE: THIS COURSE

```
44 49 \begin{frame}{Comparison of Languages}
```

```
45 50 \begin{itemize}
```

```
46 - \item Aruoba and Fernández-Villaverde (2014), "A Comparison of Programming Languages in Economics":
```

```
51 + \item<+> Aruoba and Fernández-Villaverde (2014), "A Comparison of Programming Languages in Economics":
```

```
47 52 \end{itemize}
```

```
48 - \centering\includegraphics[scale=0.4]{languageSpeeds.png}
```

```
53 + \centering\includegraphics[scale=0.35]{languageSpeeds.png}
```

```
49 54 \begin{itemize}
```

```
50 - \item Running time usually less important development time.
```

```
51 - \item Matlab and Julia run slower, but much faster to code and debug.
```

```
55 + \item<+> Running time usually less important than development time.
```

```
56 + \item<+> MATLAB and Julia run slower, but much faster to code and debug.
```

```
57 + \item<+> Julia considerably faster than MATLAB
```

```
52 58 \end{itemize}
```

```
53 59 \end{frame}
```

## GIT RESOURCES

- [Pro Git](#): free comprehensive guide
- Git is a command line tool, but many GUIs available
  - e.g. GitKraken: normally paid, but free with GitHub Student account (can get with UR email)

## GIT RESOURCES

- [Pro Git](#): free comprehensive guide
- Git is a command line tool, but many GUIs available
  - e.g. GitKraken: normally paid, but free with GitHub Student account (can get with UR email)

## GIT RESOURCES

- [Pro Git](#): free comprehensive guide
- Git is a command line tool, but many GUIs available
  - e.g. GitKraken: normally paid, but free with GitHub Student account (can get with UR email)

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!



## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## GENERAL ADVICE

- Maintainable code
  - write your code so that it could be read by anyone: then your future self will definitely be able to read it!
  - self-explanatory names
  - ample comments
- Reproducible code
  - make sure your results can be obtained by running a single file on any computer.
  - make sure you can reproduce results obtained in the past (VCS!)
- Acquire technical skills and learn tools now—won't have time after the PhD!

## JULIA RECOMMENDATIONS

---

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`



## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`

## EXPLICIT TYPES

- Explicitly state the type when declaring fields or arguments
  - field in a struct: `V::Array{Float64, 2}`, not `V`
  - `function y(x::Float64, z::Int64)`, not `function y(x, z)`
  - faster in some cases: compiler optimizes code for specified type
  - easier to catch passing a wrong variable, e.g. grid index (`Int64`) vs grid value (`Float64`)
- Preallocate arrays with explicit types
  - e.g. need to create an array element-by-element (grid, value function, etc)
  - specify array length and element type: preallocates the exact amount of memory needed
  - `policy_fn = Array{Float64}(undef, n_z, n_k)`



## TYPE STABILITY

- Keep types stable: variables shouldn't change types
  - `x = 1; x = 1.5` is slow
  - `x = 1.0; x = 1.5` is fast
- Return of a function shouldn't change type depending on input value
  - e.g. `some_fn(x :: Float)` returns `x` if `x > 0`, else `0`
  - faster if it returns `0.0`: compiler will optimize the function for floats

## TYPE STABILITY

- Keep types stable: variables shouldn't change types
  - `x = 1; x = 1.5` is slow
  - `x = 1.0; x = 1.5` is fast
- Return of a function shouldn't change type depending on input value
  - e.g. `some_fn(x :: Float)` returns `x` if `x > 0`, else `0`
  - faster if it returns `0.0`: compiler will optimize the function for floats

## TYPE STABILITY

- Keep types stable: variables shouldn't change types
  - `x = 1; x = 1.5` is slow
  - `x = 1.0; x = 1.5` is fast
- Return of a function shouldn't change type depending on input value
  - e.g. `some_fn(x :: Float)` returns `x` if `x > 0`, else `0`
  - faster if it returns `0.0`: compiler will optimize the function for floats

## TYPE STABILITY

- Keep types stable: variables shouldn't change types
  - `x = 1; x = 1.5` is slow
  - `x = 1.0; x = 1.5` is fast
- Return of a function shouldn't change type depending on input value
  - e.g. `some_fn(x :: Float)` returns `x` if `x > 0`, else `0`
  - faster if it returns `0.0`: compiler will optimize the function for floats

## TYPE STABILITY

- Keep types stable: variables shouldn't change types
  - `x = 1; x = 1.5` is slow
  - `x = 1.0; x = 1.5` is fast
- Return of a function shouldn't change type depending on input value
  - e.g. `some_fn(x::Float)` returns `x` if `x>0`, else `0`
  - faster if it returns `0.0`: compiler will optimize the function for floats

## TYPE STABILITY

- Keep types stable: variables shouldn't change types
  - `x = 1; x = 1.5` is slow
  - `x = 1.0; x = 1.5` is fast
- Return of a function shouldn't change type depending on input value
  - e.g. `some_fn(x::Float)` returns `x` if `x>0`, else `0`
  - faster if it returns `0.0`: compiler will optimize the function for floats

## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want  $a * b + \text{scalar\_fn}(c)$  done for 10 different  $a, b, c$ 
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - @. macro: appends . everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```

## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want  $a * b + \text{scalar\_fn}(c)$  done for 10 different  $a, b, c$ 
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - @. macro: appends . everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```



## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want  $a * b + \text{scalar\_fn}(c)$  done for 10 different  $a, b, c$ 
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - @. macro: appends . everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```

## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want  $a * b + \text{scalar\_fn}(c)$  done for 10 different  $a, b, c$ 
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - @. macro: appends . everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```

## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want  $a * b + \text{scalar\_fn}(c)$  done for 10 different  $a, b, c$ 
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - @. macro: appends . everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```

## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want `a * b + scalar_fn(c)` done for 10 different a, b, c
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - @. macro: appends . everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```

## VECTORIZATION

- apply a function to each element of a vector: can write a loop or use vectorized operations
  - automatically parallelized
  - not much performance difference otherwise (although loops MUCH slower in MATLAB)
  - but vectorization often more concise and readable
- e.g. want `a * b + scalar_fn(c)` done for 10 different a, b, c
  - dot notation: append a dot after any scalar operand or function

```
vec_sum = Array{Float64}(undef, 10)
vec_sum = vec_a .* vec_b .+ scalar_fn.(vec_c)
```
  - `@.` macro: appends `.` everywhere in the line

```
vec_sum = Array{Float64}(undef, 10)
@. vec_sum = vec_a * vec_b + scalar_fn(vec_c)
```

## EXPOSITION OF ADVANCED TECHNIQUES

---

## KRUSSEL AND SMITH '98

- Aiyagari + transition dynamics
  - response of the economy to shocks
  - agents need to forecast the LOM of asset distribution
- Krusell and Smith: approximate the distribution with a sequence of moments
- Forecast error from looking at mean ( $K$ ) alone is tiny
  - $\implies$  agents need to just forecast  $K$ , not the whole distribution
  - makes solving a heterogenous agent model with aggregate shocks feasible

## KRUSSEL AND SMITH '98

- Aiyagari + transition dynamics
  - response of the economy to shocks
  - agents need to forecast the LOM of asset distribution
- Krusell and Smith: approximate the distribution with a sequence of moments
- Forecast error from looking at mean ( $K$ ) alone is tiny
  - $\implies$  agents need to just forecast  $K$ , not the whole distribution
  - makes solving a heterogeneous agent model with aggregate shocks feasible



## KRUSSEL AND SMITH '98

- Aiyagari + transition dynamics
  - response of the economy to shocks
  - agents need to forecast the LOM of asset distribution
- Krusell and Smith: approximate the distribution with a sequence of moments
- Forecast error from looking at mean ( $K$ ) alone is tiny
  - $\implies$  agents need to just forecast  $K$ , not the whole distribution
  - makes solving a heterogenous agent model with aggregate shocks feasible

## KRUSSEL AND SMITH '98

- Aiyagari + transition dynamics
  - response of the economy to shocks
  - agents need to forecast the LOM of asset distribution
- Krusell and Smith: approximate the distribution with a sequence of moments
- Forecast error from looking at mean ( $K$ ) alone is tiny
  - $\implies$  agents need to just forecast  $K$ , not the whole distribution
  - makes solving a heterogenous agent model with aggregate shocks feasible

## KRUSSEL AND SMITH '98

- Aiyagari + transition dynamics
  - response of the economy to shocks
  - agents need to forecast the LOM of asset distribution
- Krusell and Smith: approximate the distribution with a sequence of moments
- Forecast error from looking at mean ( $K$ ) alone is tiny
  - $\implies$  agents need to just forecast  $K$ , not the whole distribution
  - makes solving a heterogeneous agent model with aggregate shocks feasible

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.



## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

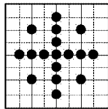
- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## PROJECTION METHODS

- To solve a model, need to find policy functions that satisfy FOCs
- Replace policy functions with parameterized approximations
  - e.g. Chebyshev polynomials
- Solve for parameters that approximately solve FOCs
  - minimize sum of squared errors over state space
- Less subject to curse of dimensionality (like log-linearization)
- Global solution (like VFI)
- Resources
  - Heer and Maußner, ch. 6
  - Fernández-Villaverde et al (2016), Solution and Estimation Methods for DSGE Models (handbook of macro chapter), section 5.

## SPARSE GRIDS

- Sparse Grids: a more efficient way of spacing grids in multi-dimensional problems
  - when function approximated is reasonably smooth, accuracy loss is small
  - objective function in VFI usually is!
- A sparse version of a 9x9 grid: 17 grid points instead of 81

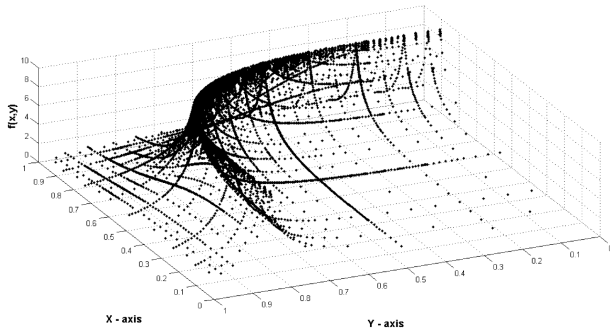


- Gets dramatic for more dimensions:

| Dimension $d$ | Full Grid $ V_4 $    | Sparse Grid $ V_4^S $ |
|---------------|----------------------|-----------------------|
| 1             | 15                   | 15                    |
| 2             | 225                  | 49                    |
| 3             | 3,375                | 111                   |
| 4             | 50,625               | 209                   |
| 5             | 759,375              | 351                   |
| 10            | $5.77 \cdot 10^{11}$ | 2,001                 |
| 20            | $3.33 \cdot 10^{23}$ | 13,201                |
| 50            | $6.38 \cdot 10^{58}$ | 182,001               |
| 100           | >Googol              | 1,394,001             |

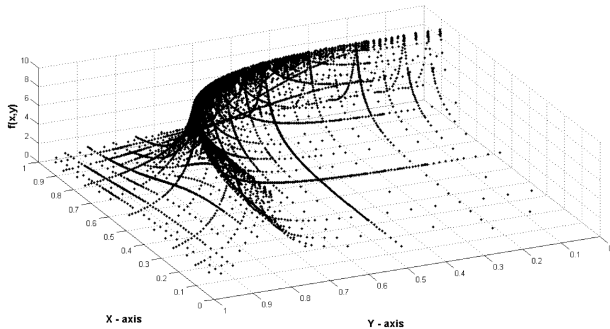
## ADAPTIVE SPARSE GRIDS

- Spacing of grids can be varied while solving the model
  - put more grid points where curvature seems higher
- Can put very few grid points in regions with relatively lower curvature
- As fast as classical sparse grids, but far more accuracy



## ADAPTIVE SPARSE GRIDS

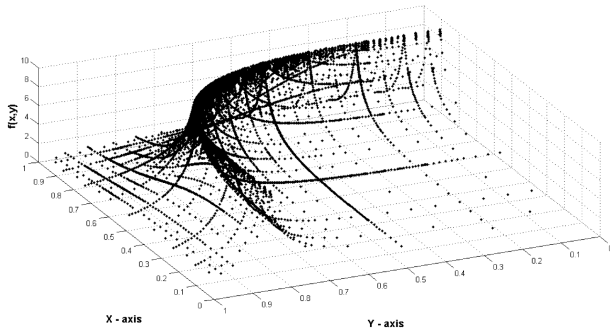
- Spacing of grids can be varied while solving the model
  - put more grid points where curvature seems higher
- Can put very few grid points in regions with relatively lower curvature
- As fast as classical sparse grids, but far more accuracy





## ADAPTIVE SPARSE GRIDS

- Spacing of grids can be varied while solving the model
  - put more grid points where curvature seems higher
- Can put very few grid points in regions with relatively lower curvature
- As fast as classical sparse grids, but far more accuracy



# ADAPTIVE SPARSE GRIDS

- Can make VFI feasible for very high-dimensional problems
- Resources
  - Brumm and Scheidegger (2017), Using Adaptive Sparse Grids To Solve High-Dimensional Dynamic Models
  - Tasmanian.jl for Julia

# ADAPTIVE SPARSE GRIDS

- Can make VFI feasible for very high-dimensional problems
- Resources
  - Brumm and Scheidegger (2017), Using Adaptive Sparse Grids To Solve High-Dimensional Dynamic Models
  - Tasmanian.jl for Julia

## ADAPTIVE SPARSE GRIDS

- Can make VFI feasible for very high-dimensional problems
- Resources
  - Brumm and Scheidegger (2017), Using Adaptive Sparse Grids To Solve High-Dimensional Dynamic Models
  - `Tasmanian.jl` for Julia

## ADAPTIVE SPARSE GRIDS

- Can make VFI feasible for very high-dimensional problems
- Resources
  - Brumm and Scheidegger (2017), Using Adaptive Sparse Grids To Solve High-Dimensional Dynamic Models
  - Tasmanian.jl for Julia

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)



## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)

## COMPUTATION ON THE GPU

- GPUs have thousands of small cores designed for rendering graphics
- CUDA is Nvidia's library that allows using GPU for general processing
- Thousands of cores  $\implies$  need a problem with thousands of parallel tasks
- Data transfer between RAM and GPU is slow  $\implies$  each task needs to be big
- VFI with a huge state space is well-suited for GPUs
  - many parallel tasks
  - more computation than communication
  - VFI can be significantly faster on your GPU than your CPU
- Resources
  - Aldrich et al. (2010), Tapping the supercomputer under your desk: Solving dynamic equilibrium models with graphics processors
  - CUDA.jl package for Julia (and its documentation)